

The Garage Problem: Real Estate Recommendations for Kings County California



By Jordan Loewen-Colón October 14th 2022

The Business Problem

King's County Realtors are interested in whether or not they should renovate homes before trying to sell. Specifically, they'd like to know how much adding a garage might affect price, and if so, what size of garage.

Recommendations:

Based on our models and analysis, we recommend that if renovations are going to occur, its best to target square footage of living space, but if renovations are going to include the garage, it is probably worth focusing on homes that have no garage and adding a 1-car sized garage, rather than increasing the size of an existing garage.

Step 1: Data Understanding

We will analyze the 2022 data from King's Country to try and offer accurate recommendations.

We begin by importing the proper tools and then the data itself.

```
In [1]: # Import Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.formula.api as smf
import statsmodels.api as sm
from scipy import stats
from sklearn.linear_model import LinearRegression

import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
warnings.filterwarnings("ignore")

#Import data
kcdf=pd.read_csv(r"C:\Users\legac\data\kc_house_data.csv")
```

```
In [2]: #view data
print(kcdf.shape)
kcdf.info()
```

```
(30155, 25)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30155 entries, 0 to 30154
Data columns (total 25 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    30155 non-null  int64
1   date                 30155 non-null  object
2   price               30155 non-null  float64
3   bedrooms            30155 non-null  int64
4   bathrooms           30155 non-null  float64
5   sqft_living         30155 non-null  int64
6   sqft_lot            30155 non-null  int64
7   floors              30155 non-null  float64
8   waterfront          30155 non-null  object
9   greenbelt           30155 non-null  object
10  nuisance             30155 non-null  object
11  view                 30155 non-null  object
12  condition            30155 non-null  object
13  grade                30155 non-null  object
14  heat_source          30123 non-null  object
15  sewer_system         30141 non-null  object
16  sqft_above           30155 non-null  int64
17  sqft_basement        30155 non-null  int64
18  sqft_garage          30155 non-null  int64
19  sqft_patio           30155 non-null  int64
20  yr_built             30155 non-null  int64
21  yr_renovated         30155 non-null  int64
22  address              30155 non-null  object
23  lat                  30155 non-null  float64
24  long                 30155 non-null  float64
dtypes: float64(5), int64(10), object(10)
memory usage: 5.8+ MB
```

The dataset has 30155 entries and 25 columns with a mix of string values, floats, and integers. Bathrooms as float makes sense, but "floors" as float seems odd. It is not clear what elements are contained in some of the object categories like "grade" or "nuisance." It also looks like we have some missing entries for "sewer_system" and "heat_source." Let's take a closer look at the data entries.

```
In [3]: kcdf.head()
```

Out[3]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	greenbelt	...	sewer_system	sqft_above	sqft_basement	sqft_garage
0	7399300360	5/24/2022	675000.0	4	1.0	1180	7140	1.0	NO	NO	...	PUBLIC	1180		0
1	8910500230	12/13/2021	920000.0	5	2.5	2770	6703	1.0	NO	NO	...	PUBLIC	1570		1570
2	1180000275	9/29/2021	311000.0	6	2.0	2880	6156	1.0	NO	NO	...	PUBLIC	1580		1580
3	1604601802	12/14/2021	775000.0	3	3.0	2160	1400	2.0	NO	NO	...	PUBLIC	1090		1070
4	8562780790	8/24/2021	592500.0	2	2.0	1120	758	2.0	NO	NO	...	PUBLIC	1120		550



It looks like there are some houses with no garages or which have never been renovated. It's also not clear how some of these columns are going to help us answer our problem, so we will probably end up focusing primarily on "yr_renovated," and the "sqft" categories. Now that we have some idea of what the values of the columns look like, let's sort by values (specifically "price") to see if we notice any outliers.

```
In [4]: kcdf.describe().sort_values("price")
```

Out[4]:

	id	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	sqft_above	sqft_basement	sqft_garage	sqft
min	1.000055e+06	2.736000e+04	0.000000	0.000000	3.000000	4.020000e+02	1.000000	2.000000	0.000000	0.000000	0.0
count	3.015500e+04	3.015500e+04	30155.000000	30155.000000	30155.000000	3.015500e+04	30155.000000	30155.000000	30155.000000	30155.000000	30155.0
25%	2.064175e+09	6.480000e+05	3.000000	2.000000	1420.000000	4.850000e+03	1.000000	1180.000000	0.000000	0.000000	40.0
50%	3.874011e+09	8.600000e+05	3.000000	2.500000	1920.000000	7.480000e+03	1.500000	1560.000000	0.000000	400.000000	150.0
std	2.882587e+09	8.963857e+05	0.981612	0.889556	974.044318	6.038260e+04	0.567717	878.306131	579.631302	285.770536	245.3
mean	4.538104e+09	1.108536e+06	3.413530	2.334737	2112.424739	1.672360e+04	1.543492	1809.826098	476.039396	330.211142	217.4
75%	7.287100e+09	1.300000e+06	4.000000	3.000000	2619.500000	1.057900e+04	2.000000	2270.000000	940.000000	510.000000	320.0
max	9.904000e+09	3.075000e+07	13.000000	10.500000	15360.000000	3.253932e+06	4.000000	12660.000000	8020.000000	3580.000000	4370.0

Some odd things to notice here is that there seem to be houses without bedrooms or bathrooms. It also looks like there might be some outliers on the larger end as well, with houses containing 13 bedrooms and 10.5 bathrooms. Also looks like there might be a house with only 3sqft of living space, so we are going to have to fix that. We're also going to need to separate out the houses that already have garages from those that do not. But before that, let's check for duplicates.

```
In [5]: # Check for duplicates
kcdf.duplicated(subset = ["id", "date"]).sum() == 0
```

Out[5]: False

So this is going to be a problem. We will need to remove them.

```
In [6]: #Checking Outliers
print(kcdf['bedrooms'].value_counts())
print(kcdf['bathrooms'].value_counts())
```

```
3      12754
4      9597
2      3936
5      2798
6       498
1       391
7        80
0        44
8        38
9        14
10         3
13         1
11         1
Name: bedrooms, dtype: int64
2.5      8475
2.0     7349
1.0     4576
3.0     4117
3.5     2266
1.5     1808
4.0      645
4.5      533
5.0      145
5.5      104
6.0       45
0.0       31
6.5       25
7.5       12
7.0       12
0.5        5
9.5        2
8.0        2
8.5        1
10.0       1
10.5       1
Name: bathrooms, dtype: int64
```


It's also still not clear how useful some of the other categories will be for getting an accurate answer to our problem. Specifically with "waterfront", "greenbelt", "nuisance", "condition", and "grade."

```
In [7]: #Checking Values of Categories
print(kcdf['waterfront'].value_counts())
print(kcdf['greenbelt'].value_counts())
print(kcdf['nuisance'].value_counts())
print(kcdf['condition'].value_counts())
print(kcdf['grade'].value_counts())
```

```
NO      29636
YES       519
Name: waterfront, dtype: int64
NO     29382
YES      773
Name: greenbelt, dtype: int64
NO     24893
YES     5262
Name: nuisance, dtype: int64
Average    18547
Good       8054
Very Good  3259
Fair       230
Poor        65
Name: condition, dtype: int64
7 Average    11697
8 Good       9410
9 Better     3806
6 Low Average 2858
10 Very Good 1371
11 Excellent  406
5 Fair       393
12 Luxury    122
4 Low        51
13 Mansion   24
3 Poor       13
2 Substandard 2
1 Cabin       2
Name: grade, dtype: int64
```

With such small values in the "waterfront", "greenbelt", and "nuisance" categories, those are probably ripe for dropping. It looks like if we dropped some of the outliers for "grade" it might end up being useful. And "condition" could be worth keeping to see how it affects our renovation coefficients. We can double check our intuition by visualizing some of the data already.

```
In [8]: #Checking For multicollinearity
#Pearson Correlation
sns.heatmap(kcdf.corr(),annot=True,cmap='RdYlGn',linewidths=0.2)
fig=plt.gcf()
fig.set_size_inches(30,25)
plt.show()
```



According to the our Pearson corellation calculations, our intuitions look right. We can drop quite a few of these categories that seem to have minimal relevance to our "price."

Step 2: Data Preperation

- In preparing the data, we will primarily focus on elements we think will affect our numbers involving renovations broadly, and garage additions more specifically. Since we are dealing with both continuous numbers and integers, we may end up needing some log transformations. We will test our model by looking at the correlation between price per sqft of house. According to according to the website www.fixr.com (<http://www.fixr.com>), the cost of building a home in California is roughly "400 and 600 per square foot." So that will be a good target for checking the accuracy of our data prep.

```
In [9]: #Create a copy of DataFrame for preperation
df = kcdf.copy(deep=True)
```

```
In [10]: #Dropping Duplicates
df.drop_duplicates(subset=['id', 'date'])

#Dropping Columns
df.drop(['id', 'date', 'view', 'lat', 'long', 'floors', 'address', 'sqft_above', 'sqft_basement', 'waterfront', 'greenbelt', 'sewage'])

#No need to drop "NaNs" because we won't be using the sewer or heating columns.
df.info()
```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30155 entries, 0 to 30154
Data columns (total 12 columns):
Column Non-Null Count Dtype
--- ---
0 price 30155 non-null float64
1 bedrooms 30155 non-null int64
2 bathrooms 30155 non-null float64
3 sqft_living 30155 non-null int64
4 sqft_lot 30155 non-null int64
5 nuisance 30155 non-null object
6 condition 30155 non-null object
7 grade 30155 non-null object
8 sqft_garage 30155 non-null int64
9 sqft_patio 30155 non-null int64
10 yr_built 30155 non-null int64
11 yr_renovated 30155 non-null int64
dtypes: float64(2), int64(7), object(3)
memory usage: 2.8+ MB

Visualizing new dataframe for correlations

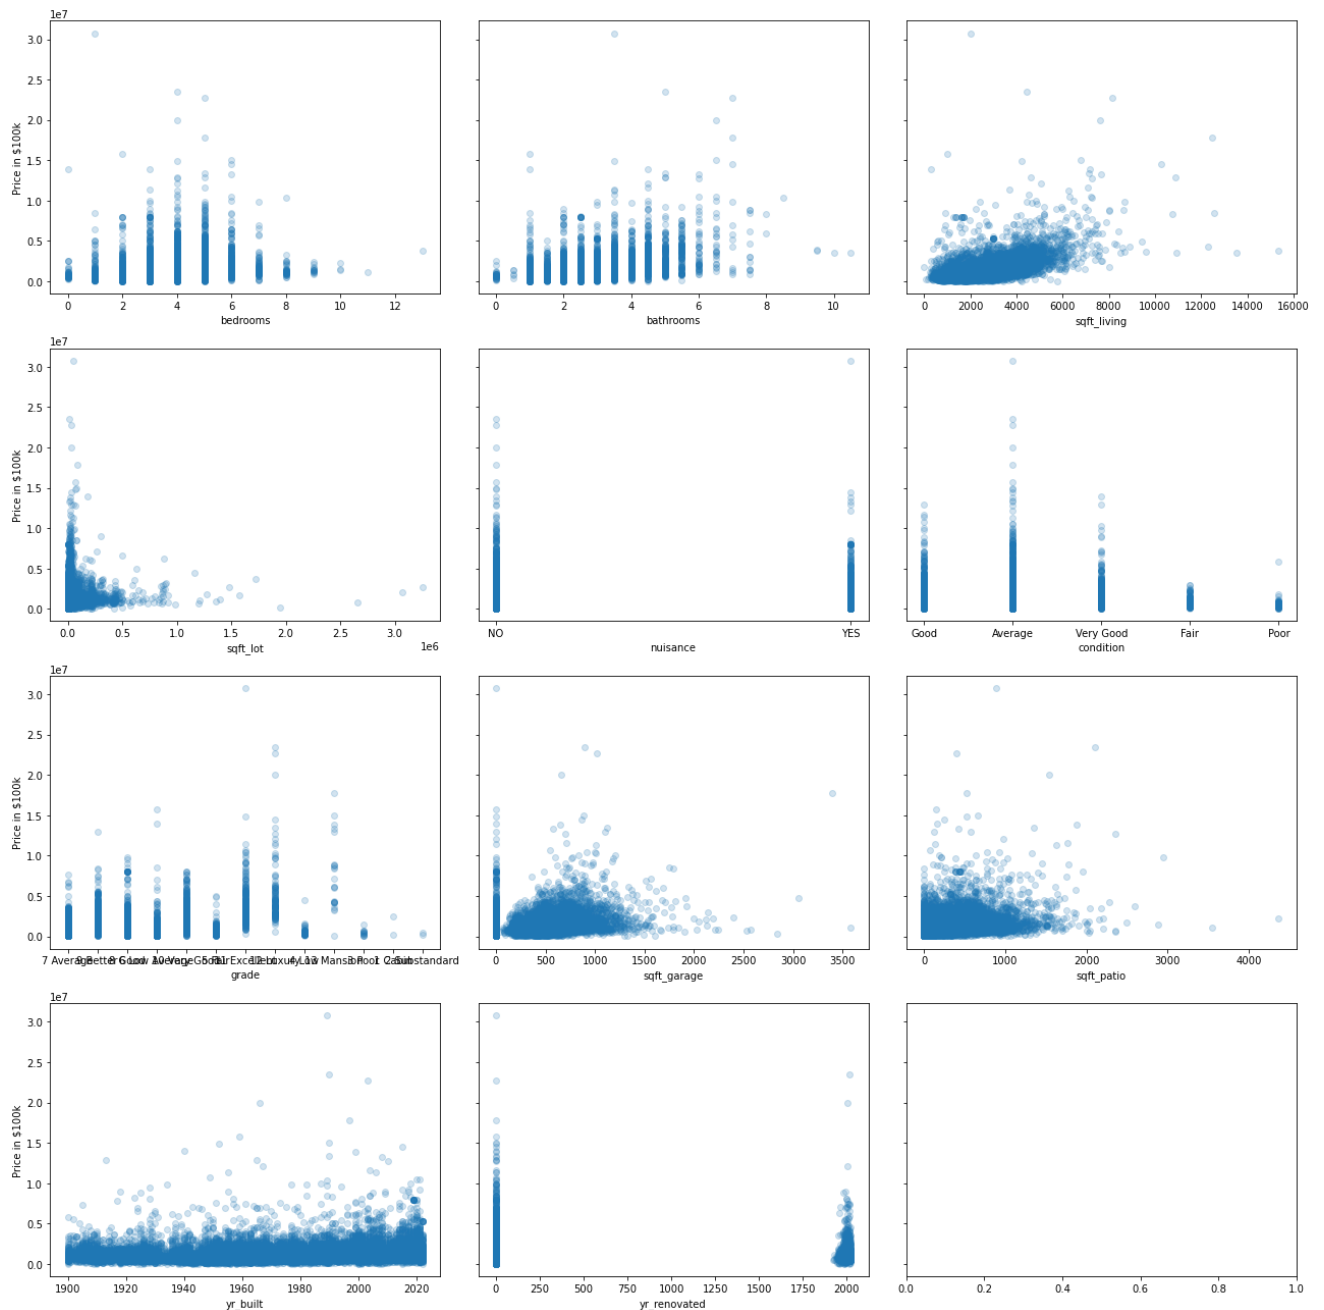
```
In [11]: #Setting X and Y
y = df["price"]
x = df.drop("price", axis=1)

fig, axes = plt.subplots(nrows=4, ncols=3, figsize=(18,18), sharey=True)

for i, column in enumerate(x.columns):
    # Locate applicable axes
    row = i // 3
    col = i % 3
    ax = axes[row][col]

    # Plot feature vs. y and label axes
    ax.scatter(x[column], y, alpha=0.2)
    ax.set_xlabel(column)
    if col == 0:
        ax.set_ylabel("Price in $100k")

fig.tight_layout()
```




Looks like we will need to fix the outliers in bedrooms, bathrooms, and our sqft columns. The sqft_lot also seems really odd. We would expect price to go up the larger the lot size, but this plot seems to say the opposite. It also looks like we can adjust the "condition" and "grade" categories into ordinal numbers so that our models will learn the proper relationship to the data. We will also have to adjust "yr_renovated" to make it more useful, and make a category frame for whether a home has a garage or not.

```
In [12]: print('Shape before filtering', df.shape)

#Let's try converting sqft_lot into 1/x to see if that at least helps give us a better visual
df['inv_sqft_lot'] = df['sqft_lot'].apply(lambda x: 1/x)

sns.scatterplot(data=df, x='inv_sqft_lot', y='price');
```

Shape before filtering (30155, 12)



That definitely looks more linear! But it is no guarantee that the fix will affect the model.

Now to fix "grade" and "condition." Turning them from strings to integers allows our model to see the relationship between the numbers more clearly and hopefully will give us more accuracy.

```
In [13]: #Convert str to int for 'condition'
df['condition'] = df['condition'].replace(['Poor', 'Fair', 'Average', 'Good', 'Very Good'], [0, 1, 2, 3, 4])

#Convert str to int for 'grade'
df['grade'] = df['grade'].replace(['1 Cabin', '2 Substandard', '3 Poor', '4 Low', '5 Fair',
                                   '6 Low Average', '7 Average', '8 Good', '9 Better',
                                   '10 Very Good', '11 Excellent', '12 Luxury', '13 Mansion'],
                                   [0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12])
```

Outliers

```
In [14]: #Deal with easy Outliers

print('Shape before filtering', df.shape)

#Remove "mansion" level houses and foreclosures (based on foreclosure listings)
df.drop(df[(df['price'] >= 5000000) | (df['price'] <= 300000)].index, inplace = True)

#A House must have a bedroom and a bathroom. Mansions are such an outlier that our model won't make much use of them.
df.drop(df[(df['bedrooms'] == 0) | (df['bedrooms'] >= 10)].index, inplace = True)
df.drop(df[(df['bathrooms'] == 0) | (df['bathrooms'] >= 8)].index, inplace = True)

#Drop anything well below "small home" and "mansion" sized in sqft_living
df.drop(df[(df['sqft_living'] <= 500) | (df['sqft_living'] >= 7000)].index, inplace = True)

#Use Z-score to drop anything above three standard deviations for sqft_lot
df['sqft_lot'] = df['sqft_lot'].loc[(np.abs(stats.zscore(df['sqft_lot'])) < 3)]
df = df.dropna(subset = ['sqft_lot'])

#Drop massive sized garages
df.drop(df[(df['sqft_garage'] >= 800)].index, inplace = True)

print('Shape after filtering', df.shape)
```

Shape before filtering (30155, 13)

Shape after filtering (27642, 13)

So after filtering, we lost about 2513 rows. We probably still have room to finetune the data by doing some logarithmic transformations, normalization, and removing outliers within three standard deviations. But we will save that for after we've done some baseline modeling below.

Next! Since we are specifically looking to answer the question about "rennovations" we will prepare our data by creating a column that determines whether or not a home has been renovated at all.


```
In [15]: df['renovated']=''
def determine_reno(df):
    t=''
    x=df['yr_renovated']
    if(x==0):
        t=0
    elif(x>0):
        t=1
    return t

df['renovated']=df.apply(determine_reno,axis=1)
df.renovated.value_counts()
```

```
Out[15]: 0    26410
         1     1232
         Name: renovated, dtype: int64
```

Since there aren't actually that many renovated houses in our data set, we won't bother trying to see how "recency" might affect the price.

Next, we will do a similar adjustment with our garage data. First checking to see which houses actually have a garage, and then checking to see how big. According to [ShedsUnlimited \(https://shedsunlimited.net/blog/how-large-is-a-one-car-garage/#:~:text=The%20minimum%20size%20of%20a,garage%20to%20fit%20inside%20comfortably,\)](https://shedsunlimited.net/blog/how-large-is-a-one-car-garage/#:~:text=The%20minimum%20size%20of%20a,garage%20to%20fit%20inside%20comfortably,) the minimum size for a garage is 180ft, but really should be about 240ft in order to actually fit a car.

```
In [16]: #Create a function to separate homes with a garage from those without
df['garage']=''
def determine_garage(df):
    t=''
    x=df['sqft_garage']
    if(x<=239):
        t="No"
    elif(x>=240):
        t="Yes"
    return t

df['garage']=df.apply(determine_garage,axis=1)
df.garage.value_counts()
```

```
Out[16]: Yes    16814
         No     10828
         Name: garage, dtype: int64
```

```
In [17]: #Create a function to categorize by size of garage
df['garage_size']=''
def garage_sizer(df):
    t=''
    x=df['sqft_garage']
    if(x<=239):
        t=0
    elif(x>=240 and x<=359):
        t=1
    elif(x>=360 and x<=704):
        t=2
    elif(x>704):
        t=3
    return t

df['garage_size']=df.apply(garage_sizer,axis=1)
df.garage_size.value_counts()
```

```
Out[17]: 2    13172
         0    10828
         1     2610
         3     1032
         Name: garage_size, dtype: int64
```

With our garage size categories, it looks like we have a nice set of numbers to work with on our models. We may still need to adjust for some outliers, but overall, it makes sense that there seem to be more "medium" or two car garages than the others.

```
In [18]: garage = df[df['garage'] == 1]
no_garage = df[df['garage'] == 0]

alpha = 0.05
garage_p_val = stats.ttest_ind(garage.price, no_garage.price, equal_var=False)[1]
print("Garage vs No Garage T-test P Value: ", garage_p_val)
if garage_p_val < 0.05:
    print("Having a garage vs not having a garage is statistically relevant to average property value")
else:
    print("accept null hypothesis")
```

Garage vs No Garage T-test P Value: nan
accept null hypothesis

```
In [19]: df.describe().sort_values("price")
```

Out[19]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	condition	grade	sqft_garage	sqft_patio	yr_built	yr_reno
count	2.764200e+04	27642.000000	27642.000000	27642.000000	27642.000000	27642.000000	27642.000000	27642.000000	27642.000000	27642.000000	27642.000000
min	3.020000e+05	1.000000	0.500000	510.000000	402.000000	0.000000	1.000000	0.000000	0.000000	1900.000000	0.000000
std	6.082013e+05	0.930925	0.824200	840.587017	15503.604051	0.705118	1.027392	246.656641	221.281240	32.355725	412.400000
25%	6.500000e+05	3.000000	2.000000	1400.000000	4685.250000	2.000000	6.000000	0.000000	40.000000	1952.000000	0.000000
50%	8.500000e+05	3.000000	2.500000	1880.000000	7236.000000	2.000000	6.000000	380.000000	140.000000	1976.000000	0.000000
mean	1.041877e+06	3.391035	2.292725	2031.490449	10399.715035	2.483395	6.581253	303.278489	203.885826	1974.609363	89.000000
75%	1.251750e+06	4.000000	2.500000	2520.000000	9900.000000	3.000000	7.000000	490.000000	300.000000	2003.000000	0.000000
max	4.995000e+06	7.000000	9.500000	6860.000000	192212.000000	4.000000	12.000000	790.000000	2880.000000	2022.000000	2022.000000

Data Modeling

Now with our data prepared and in hand, it's time to create some models. First, we will check to see the accuracy of our data preparation by checking price per sqft.

```
In [20]: #Set "price" as y
y = df['price']
X = df['sqft_living']

model = sm.OLS(y, sm.add_constant(X))
results = model.fit()

print(results.summary())
```

```
OLS Regression Results
=====
Dep. Variable:      price      R-squared:      0.397
Model:              OLS      Adj. R-squared:    0.397
Method:             Least Squares      F-statistic:    1.818e+04
Date:               Thu, 12 Jan 2023      Prob (F-statistic): 0.00
Time:               18:59:11      Log-Likelihood:   -4.0038e+05
No. Observations:   27642      AIC:              8.008e+05
Df Residuals:       27640      BIC:              8.008e+05
Df Model:           1
Covariance Type:    nonrobust
=====
               coef      std err          t      P>|t|      [0.025      0.975]
-----
const          1.16e+05    7431.206     15.607     0.000     1.01e+05     1.31e+05
sqft_living    455.7727      3.380     134.841     0.000     449.148     462.398
=====
Omnibus:           9877.343      Durbin-Watson:      1.971
Prob(Omnibus):     0.000      Jarque-Bera (JB):    59218.875
Skew:              1.598      Prob(JB):            0.00
Kurtosis:          9.419      Cond. No.            5.75e+03
=====
```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 5.75e+03. This might indicate that there are strong multicollinearity or other numerical problems.

So it looks like the coefficient is within the right range, but the R-squared is very low. We will need to adjust in order to feel more confident in our model's predictions. Let's see if our log transformation gives us a better value:

```
In [21]: # Log-transforming chosen variables
df["log_sqft_living"] = np.log(df[["sqft_living"]])
```

```
In [22]: #Set X
X = df["log_sqft_living"]

model = sm.OLS(y, sm.add_constant(X))
results = model.fit()

print(results.summary())
```

```

=====
                        OLS Regression Results
=====
Dep. Variable:          price      R-squared:                0.328
Model:                  OLS      Adj. R-squared:             0.328
Method:                 Least Squares    F-statistic:         1.352e+04
Date:                   Thu, 12 Jan 2023    Prob (F-statistic):    0.00
Time:                   18:59:11      Log-Likelihood:       -4.0186e+05
No. Observations:       27642      AIC:                  8.037e+05
Df Residuals:           27640      BIC:                  8.037e+05
Df Model:                1
Covariance Type:        nonrobust
=====
                        coef      std err          t      P>|t|      [0.025      0.975]
-----
const                -5.345e+06    5.5e+04    -97.158      0.000    -5.45e+06    -5.24e+06
log_sqft_living      8.478e+05    7291.580    116.269      0.000     8.33e+05     8.62e+05
=====
Omnibus:                 10779.110    Durbin-Watson:           1.975
Prob(Omnibus):           0.000    Jarque-Bera (JB):        63271.714
Skew:                    1.778    Prob(JB):                0.00
Kurtosis:                9.503    Cond. No.                141.
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Woah! Looks worse. What if we adjust using Z-score?

```
In [23]: #First seperate continuous #s
df_cont = df[['price', 'sqft_living', 'sqft_garage', 'bedrooms', 'bathrooms', 'grade', 'condition']].copy()

#Remove outliers based on Z-score. Tinkered with the # to see what gave the best model.
df_std = df[(np.abs(stats.zscore(df_cont)) < 3).all(axis=1)]

#Check to see how many outliers were removed
print(len(df)-len(df_std))
```

962

```
In [24]: #Set X and y
y = df_std["price"]
X = df_std["sqft_living"]

model = sm.OLS(y, sm.add_constant(X))
results = model.fit()

print(results.summary())
```

```
=====
                        OLS Regression Results
=====
```

Dep. Variable:	price	R-squared:	0.341
Model:	OLS	Adj. R-squared:	0.341
Method:	Least Squares	F-statistic:	1.378e+04
Date:	Thu, 12 Jan 2023	Prob (F-statistic):	0.00
Time:	18:59:11	Log-Likelihood:	-3.8129e+05
No. Observations:	26680	AIC:	7.626e+05
Df Residuals:	26678	BIC:	7.626e+05
Df Model:	1		
Covariance Type:	nonrobust		

```
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	2.504e+05	6654.472	37.630	0.000	2.37e+05	2.63e+05
sqft_living	370.0362	3.152	117.399	0.000	363.858	376.214

```
=====
```

Omnibus:	3950.587	Durbin-Watson:	2.010
Prob(Omnibus):	0.000	Jarque-Bera (JB):	7386.379
Skew:	0.942	Prob(JB):	0.00
Kurtosis:	4.760	Cond. No.	5.89e+03

```
=====
```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
 [2] The condition number is large, 5.89e+03. This might indicate that there are strong multicollinearity or other numerical problems.

This time its better than the log transformation, but it is still worse than our original test. So that means we probably did enough cleaning during our data prep stage. Before we add the renovation data and garage data, let's see which variables give us the most accurate model.

```
In [25]: #Set X and y
y = df["price"]
X = df[['sqft_living', 'sqft_lot', 'bedrooms', 'bathrooms', 'yr_built', "grade", 'renovated', 'condition', 'garage']]

#Dummies for Garage, Grade, and Condition
X = pd.get_dummies(X, columns=['garage', 'renovated', "grade", "condition"], drop_first=True)

model = sm.OLS(y, sm.add_constant(X))
results = model.fit()

print(results.summary())
```

```

=====
                        OLS Regression Results
=====
Dep. Variable:          price      R-squared:                0.533
Model:                  OLS        Adj. R-squared:            0.532
Method:                 Least Squares   F-statistic:           1431.
Date:                   Thu, 12 Jan 2023   Prob (F-statistic):    0.00
Time:                   18:59:11      Log-Likelihood:        -3.9685e+05
No. Observations:       27642          AIC:                   7.937e+05
Df Residuals:           27619          BIC:                   7.939e+05
Df Model:                22
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	8.845e+06	4.74e+05	18.681	0.000	7.92e+06	9.77e+06
sqft_living	258.1625	5.966	43.273	0.000	246.469	269.856
sqft_lot	-0.0760	0.167	-0.454	0.650	-0.404	0.252
bedrooms	-3.947e+04	3806.885	-10.368	0.000	-4.69e+04	-3.2e+04
bathrooms	8.269e+04	5178.060	15.969	0.000	7.25e+04	9.28e+04
yr_built	-4373.9691	113.673	-38.478	0.000	-4596.774	-4151.164
garage_Yes	-3.016e+04	6045.465	-4.989	0.000	-4.2e+04	-1.83e+04
renovated_1	7.174e+04	1.3e+04	5.517	0.000	4.63e+04	9.72e+04
grade_2	5.428e+05	5.88e+05	0.922	0.356	-6.1e+05	1.7e+06
grade_3	7.339e+04	4.31e+05	0.170	0.865	-7.71e+05	9.18e+05
grade_4	-1.457e+05	4.22e+05	-0.345	0.730	-9.73e+05	6.82e+05
grade_5	-1.579e+05	4.22e+05	-0.375	0.708	-9.84e+05	6.69e+05
grade_6	-4.18e+04	4.22e+05	-0.099	0.921	-8.68e+05	7.85e+05
grade_7	1.573e+05	4.22e+05	0.373	0.709	-6.69e+05	9.84e+05
grade_8	5.112e+05	4.22e+05	1.212	0.226	-3.16e+05	1.34e+06
grade_9	1.008e+06	4.22e+05	2.388	0.017	1.81e+05	1.83e+06
grade_10	1.429e+06	4.23e+05	3.379	0.001	6e+05	2.26e+06
grade_11	1.404e+06	4.29e+05	3.269	0.001	5.62e+05	2.25e+06
grade_12	1.889e+06	5.93e+05	3.187	0.001	7.27e+05	3.05e+06
condition_1	1.211e+05	7.15e+04	1.693	0.090	-1.91e+04	2.61e+05
condition_2	1.312e+05	6.53e+04	2.007	0.045	3072.456	2.59e+05
condition_3	1.525e+05	6.53e+04	2.335	0.020	2.45e+04	2.81e+05
condition_4	1.993e+05	6.56e+04	3.039	0.002	7.07e+04	3.28e+05

```

=====
Omnibus:                 10623.350   Durbin-Watson:           1.963
Prob(Omnibus):            0.000   Jarque-Bera (JB):        90241.675
Skew:                     1.619   Prob(JB):                 0.00
Kurtosis:                 11.238   Cond. No.                 1.10e+07
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.1e+07. This might indicate that there are strong multicollinearity or other numerical problems.

With this model, it looks like houses with all other factors the same, but with a garage end up being valued at about 30k less than the same home without a garage.

However, while this model explains about 53% of our data, quite a few of our variables, specifically in the "grade" section, don't quite meet our threshold $\alpha = 0.05$. Let's see the difference of garage size is taken into account.

```
In [26]: #Set X and y
y = df["price"]
X = df[['sqft_living', 'bedrooms', 'bathrooms', 'yr_built', "grade", 'condition', 'renovated', 'garage_size']]

#Dummies for Renovated
X = pd.get_dummies(X, columns=["renovated"], drop_first=True)

#Dummies for grade, and drop "7 Average" as reference category
X = pd.get_dummies(X, columns=["garage_size"])
X = X.drop("garage_size_0", axis=1)

#Dummies for grade, and drop "7 Average" as reference category
X = pd.get_dummies(X, columns=["grade"])
X = X.drop("grade_7", axis=1)

#Dummies for 'condition' and drop 'Average' as reference category
X = pd.get_dummies(X, columns=["condition"])
X = X.drop("condition_2", axis=1)

model = sm.OLS(y, sm.add_constant(X))
results = model.fit()

print(results.summary())
```

```
=====
                        OLS Regression Results
=====
```

Dep. Variable:	price	R-squared:	0.534
Model:	OLS	Adj. R-squared:	0.533
Method:	Least Squares	F-statistic:	1374.
Date:	Thu, 12 Jan 2023	Prob (F-statistic):	0.00
Time:	18:59:11	Log-Likelihood:	-3.9682e+05
No. Observations:	27642	AIC:	7.937e+05
Df Residuals:	27618	BIC:	7.939e+05
Df Model:	23		
Covariance Type:	nonrobust		

```
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	8.985e+06	2.27e+05	39.664	0.000	8.54e+06	9.43e+06
sqft_living	263.9585	5.930	44.511	0.000	252.335	275.582
bedrooms	-3.928e+04	3795.404	-10.350	0.000	-4.67e+04	-3.18e+04
bathrooms	8.194e+04	5161.863	15.874	0.000	7.18e+04	9.21e+04
yr_built	-4302.6286	114.421	-37.604	0.000	-4526.899	-4078.358
renovated_1	6.951e+04	1.3e+04	5.352	0.000	4.41e+04	9.5e+04
garage_size_1	3887.3469	9216.170	0.422	0.673	-1.42e+04	2.2e+04
garage_size_2	-3.84e+04	6517.193	-5.892	0.000	-5.12e+04	-2.56e+04
garage_size_3	-1.124e+05	1.49e+04	-7.547	0.000	-1.42e+05	-8.32e+04
grade_1	-1.641e+05	4.21e+05	-0.390	0.696	-9.89e+05	6.61e+05
grade_2	3.783e+05	4.21e+05	0.899	0.368	-4.46e+05	1.2e+06
grade_3	-8.124e+04	8.91e+04	-0.912	0.362	-2.56e+05	9.34e+04
grade_4	-2.992e+05	2.59e+04	-11.535	0.000	-3.5e+05	-2.48e+05
grade_5	-3.159e+05	1.1e+04	-28.616	0.000	-3.38e+05	-2.94e+05
grade_6	-2.006e+05	6724.100	-29.838	0.000	-2.14e+05	-1.87e+05
grade_8	3.565e+05	8902.655	40.046	0.000	3.39e+05	3.74e+05
grade_9	8.616e+05	1.57e+04	54.849	0.000	8.31e+05	8.92e+05
grade_10	1.293e+06	3.38e+04	38.234	0.000	1.23e+06	1.36e+06
grade_11	1.262e+06	8.1e+04	15.590	0.000	1.1e+06	1.42e+06
grade_12	1.704e+06	4.16e+05	4.093	0.000	8.88e+05	2.52e+06
condition_0	-1.319e+05	6.53e+04	-2.021	0.043	-2.6e+05	-3992.735
condition_1	-1.023e+04	3.02e+04	-0.338	0.735	-6.95e+04	4.9e+04
condition_3	2.275e+04	6245.103	3.642	0.000	1.05e+04	3.5e+04
condition_4	6.813e+04	8699.732	7.831	0.000	5.11e+04	8.52e+04

```
=====
```

Omnibus:	10622.554	Durbin-Watson:	1.963
Prob(Omnibus):	0.000	Jarque-Bera (JB):	90383.905
Skew:	1.618	Prob(JB):	0.00
Kurtosis:	11.246	Cond. No.	4.94e+05

```
=====
```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
 [2] The condition number is large, 4.94e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [27]: #Calculating absolute error
from sklearn.metrics import mean_absolute_error

y_pred = results.predict(sm.add_constant(X))
mean_absolute_error(y, y_pred)
```

Out[27]: 283292.8977197938

Data Understanding

Results of model: This model explains about 53.3% of the variance in our data. This model's F-statistic is statistically significant compared to our alpha of 0.05. Most of the coefficients are statistically significant when compared to our alpha of .05.

Interpretations: For a house with no garage, of average grade and condition, and with no renovations we would expect the house to be about 70k less than a home that is renovated. We expect that same house to sell for about 4k more with a 1-car garage. For each additional 1 square foot in living space size and all other features remaining the same, we would expect the house to gain about \$263.

Conclusions

According to our models, renovating a home, specifically targeting square footage of living space, will have a significant increase on the value of the home. While it looks like adding a garage can increase the value of a home in some scenarios, it is dependent on the size of garage and other factors. We would be hard pressed to recommend adding a 1-car garage to a home that does not have one, given that the price difference is only about 4k increase.

To that end, based on our models and analysis, we recommend that if renovations are going to occur, it's best to target square footage of living space, but if renovations are going to include the garage, it is probably worth focusing on homes that have no garage and adding a 1-car sized garage, rather than increasing the size of an existing garage.

Next Steps

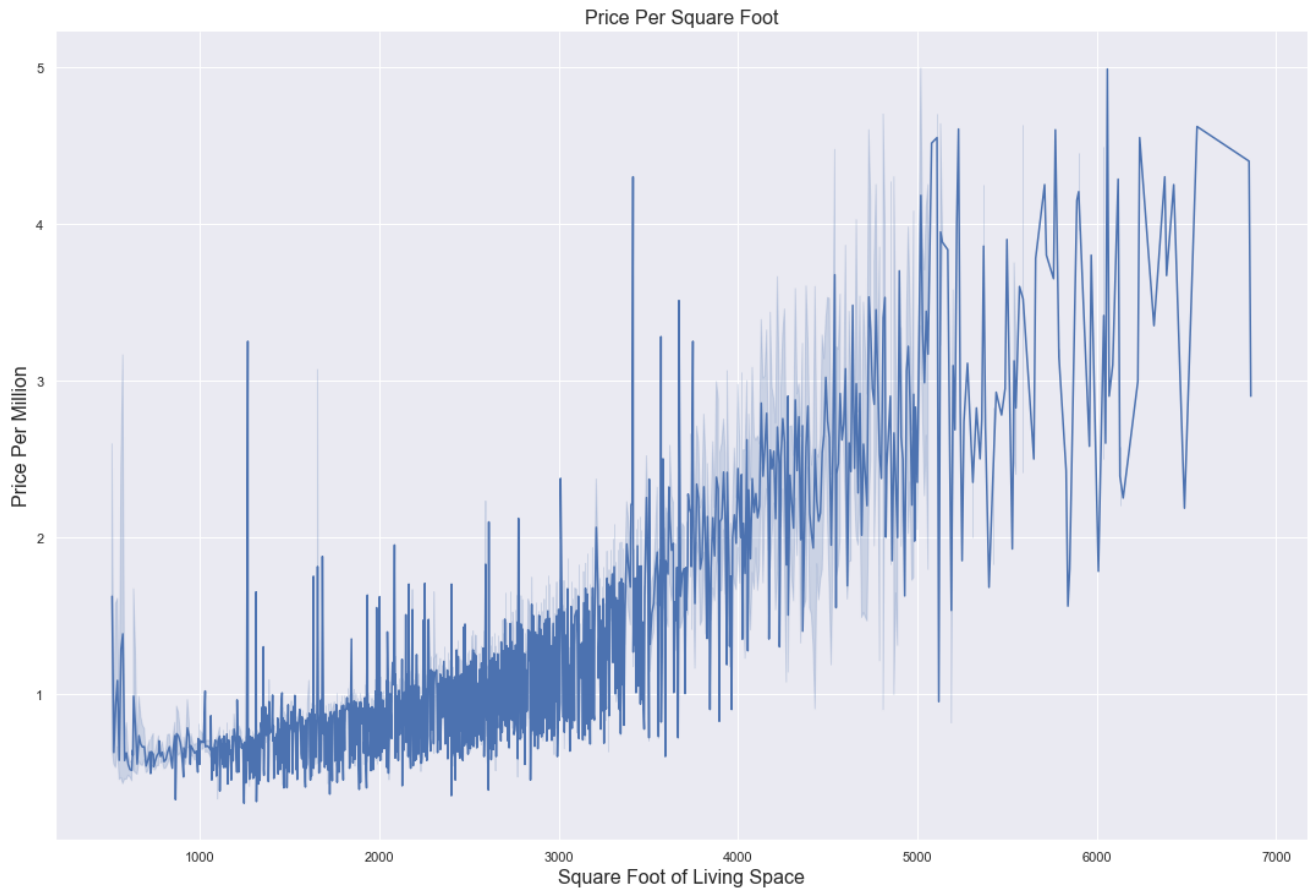
With more time, we could include other factors, like zip code, that might increase the accuracy of our model and thus update our recommendations.

Additional Visualizations


```
In [28]: #Price per Square Foot of Living Space
price = df['price']/1000000

sns.set_theme(style="darkgrid")


sns.set(rc={"figure.figsize":(20, 5)})
sns.relplot(data=df, x="sqft_living", y=price, kind="line", height=10, aspect=1.5);
plt.title('Price Per Square Foot', fontsize=16)
plt.xlabel('Square Foot of Living Space', fontsize=16)
plt.ylabel('Price Per Million', fontsize=16)
plt.ticklabel_format(style='plain', axis='x')
plt.ticklabel_format(style='plain', axis='y');
```




```
In [29]: slope, intercept, r_value, p_value, std_err = stats.linregress(df['sqft_living'], y)
sns.regplot(x="sqft_living", y=price, data=df, label='',
            line_kws={'label': "y={0:.1f}x+{1:.1f}".format(slope, intercept), "color": "red"})
plt.title('Price Per Square Foot', fontsize=16)
plt.xlabel('Square Foot of Living Space', fontsize=16)
plt.ylabel('Price Per Million', fontsize=16)
plt.ticklabel_format(style='plain', axis='x')
plt.ticklabel_format(style='plain', axis='y');
plt.legend();
```




```
In [30]: #Price With vs Without Garage
sns.relplot(data=df, x="sqft_living", y=price, hue="garage", height=10, aspect=1.5)
plt.title('Price Per Square Foot With or Without Garage', fontsize=16)
plt.xlabel('Square Foot of Living Space', fontsize=16)
plt.ylabel('Price Per Million', fontsize=16)
plt.ticklabel_format(style='plain', axis='x')
plt.ticklabel_format(style='plain', axis='y');
```



```
In [31]: #Price vs Size of Garage
sns.relplot(data=df, x="sqft_living", y=price, hue="garage_size", height=10, aspect=1.5);
plt.title('Price Per Square Foot with Garage Size', fontsize=16)
plt.xlabel('Square Foot of Living Space', fontsize=16)
plt.ylabel('Price Per Million', fontsize=16)
plt.ticklabel_format(style='plain', axis='x')
plt.ticklabel_format(style='plain', axis='y');
```



```
In [ ]:
```